

Girish Manjunatha Saraswathipura

gsaraswathipura@gmail.com | +1 (703)-400-9756 | [LinkedIn](#) | [GitHub](#) | Bloomington IN

Education

Master of Science in Data Science - ML, Data Mining, Advance Database, AI Research, Applied Algorithm Indiana University Bloomington, USA	08/2024 - 04/2026
Bachelor of Engineering in Computer Science - Mathematics, DevOps, Software Engineering, MongoDB Visvesvaraya Technological University, India	08/2018 - 06/2022

Work Experience

IU Innovates <i>Machine Learning Research Intern - Deep Learning, HPC, Fine Tune LLM, Statistics, Data Profiling, CNN, ANN, GPU, CUDA</i> • Researched how LLMs set task boundaries using system prompts, improving controlled response accuracy by 25%. • Explored LLM meta-learning to enable quick adaptation across tasks, cutting model retraining needs by 30%. • Built an AutoML tool that shrinks model search space, achieving 40% lower cost and 60% faster convergence during benchmarking. • Applied advanced statistical and approximation methods to extract metadata from 50M+ records, cutting profiling time by 45%. • Automated preprocessing and EDA workflows, reducing manual data exploration effort by 30% and accelerating model readiness. • Developed a deep learning architecture generator in TensorFlow & PyTorch for image and tabular data, improving convergence with 35% fewer iterations. • Leveraged NVIDIA A100 GPUs with CUDA and multithreaded processing to optimize training speed and computational efficiency by 2x.	06/2025 - 08/2025
Informatica Inc. <i>Modernization Associate Consultant - ETL optimization, Data Governance, Data Quality, Synapse SQL, AWS EC2, Azure, Snowflake</i> • Implemented a Python automation tool on AWS EC2 that reduced certificate generation time from 10 hrs to 3 hrs, saving \$4K+ monthly. • Consulted Fortune 500 companies to migrate 370+ ETL workflows (8% of enterprise pipelines) to Informatica Cloud, accelerating adoption in 6 months and earning an MVP award for impact on Volvo and Charles Schwab projects.	01/2024 - 07/2024
Informatica Inc. <i>Migration Factory Associate Specialist - Data Engineering, Data Warehouse, Data Wrangling, Oracle, PySpark, Hadoop</i> • Migrated 1,100+ pipelines from PowerCenter to IICS for Volvo's master database, boosting workflow efficiency by 21% and ensuring 98%+ reliability on 20TB monthly data. • Automated lineage checks for 300+ dimensions and profiled streaming data across Snowflake, Hadoop, Spark, Oracle, AWS Redshift, and Azure Synapse, completing a major migration 7 weeks ahead of schedule.	08/2022 - 12/2023

Projects

Optimizing ML Pipeline using Advanced Statistics - AutoML, Inferential Statistics, Regularization, PySpark, Mutual Information • Developed an AutoML + meta-learning tool with advanced statistical sampling ($\pm 3\%$ error) and PySpark optimization, reducing ML pipeline cost by 25% and runtime by 40%, backed by \$8K in grants.	
Traffic Sign Classification on AWS SageMaker - MLOps, Model Deployment, Neural Networks, AWS • Built and deployed a CNN (LeNet) on AWS SageMaker with automated preprocessing, achieving 95% accuracy and processing 100K+ predictions/day at 60% faster runtime and $\frac{2}{3}$ cloud cost.	
Audio Signal Classification - TensorFlow, Class Imbalance Handling, Data Augmentation, CNN, Signal Processing, Feature Extraction • Built an audio signal classification model on 800+ files (217 positive, 593 negative) using spectrogram features and CNNs in PyTorch, achieving 90%+ accuracy with class imbalance handling.	
Fine-Tuning LLM on Environmental Protection Data - Hugging Face Transformers, Tokenization & Embeddings, Transfer Learning • Collected and transformed Indiana environmental blogs and government records into Q&A datasets, enabling Hugging Face fine-tuning of LLaMA-8B with 15% faster convergence and improved model adaptability.	

Skills

Programming & Data: Python, C++, C, R, SQL, PostgreSQL, Synapse SQL, MongoDB, FAISS, NeoJS, Hadoop, PySpark, AWS S3, Snowflake
ML Engineering & MLOps: Feature Engineering, Feature Selection, Hyperparameter Tuning, Cross-Validation, Regularization, Ensemble Learning, Grid Search, Model Monitoring, Drift Detection, CI/CD for ML, Reproducibility, Versioning, MLflow, Experiment Tracking, NLP, PyTorch, Scikit-learn, TensorFlow, Transformers, RAG, LangChain, Foundation Models, Agentic AI, Anthropic Claude API
Data & Cloud Platforms: Databricks, Informatica, AWS (Bedrock, Glue, SageMaker), Docker, Kubernetes, ETL Pipelines, Data Governance
Mathematics & Statistics: Linear Algebra, Probability & Statistics, Calculus, gradients, Numerical Optimization, Information Theory

Certifications and Credentials: IBM Data Science Professional, AWS AI Practitioner, Informatica Certified Data Integration Professional, DeepLearning Certified TensorFlow Developer, Intel AI Specialization, MongoDB University M001, Informatica Cloud Application Integration
--